

Topology in Neural Machine Translation: A Topological Study of Transformers through Attention

Yuwei (Johnny) Meng

23 Dec 2025

[Link to GitHub Repo](#)

Abstract

The transformer architecture revolutionized Natural Language Processing, but understanding how transformers process language remains an active area of research, with NMT being particularly underexplored. This study applies TDA to NMT by analyzing attention maps from the NLLB model on 2,000 sentence pairs each from the WMT French-English and Chinese-English datasets. We compute persistent homology on attention-derived graphs and measure cross-language topological similarity using Wasserstein distances between persistence diagrams. For French-English, we find a statistically significant negative correlation between topological dissimilarity and average BLEU scores ($r = -0.132, p = 2.84 \times 10^{-9}$) after controlling for sentence length, indicating that preserving topological structures independently contributes to translation quality. However, the Chinese-English analysis revealed systematic model limitations that severely affect BLEU scores. The contrast between the statistically significant French-English correlation and the confounded Chinese-English results suggests that preserving topological structures may be particularly important for typologically related language pairs, where the model demonstrates more reliable performance.

1. Introduction

In 2017, a group of researchers at Google proudly announced the architecture of the transformer neural model, which revolutionized the field of Natural Language Processing (NLP) (Vaswani et al., 2017). Before then, neural NLP models mainly relied on recurrent structures, such as recurrent neural networks (RNNs) and Long Short-Term Memories (LSTMs), optionally with attention to boost performance. The transformer architecture, however, abandoned the recurrent structures in RNNs and LSTMs and solely utilized attention for language modeling, which was a novel but successful approach. Since the invention of transformer, NLP researchers have been applying this architecture to various NLP tasks, one of which is Machine Translation (MT). MT is an NLP task that takes a

sentence in a source language as input and outputs the translated sentence in a target language, and Neural Machine Translation (NMT) is a subfield of MT that specifically uses neural networks as the model for translation. According to Vaswani et al. (2017), the transformer model achieved a BLEU score of 41.0 on the 2014 Workshop on Machine Translation (WMT14) English-to-French benchmark, establishing a new state-of-the-art performance.

Despite the prominent performance of transformer, similar to other neural network architectures, the specific reasons behind its success remain largely unknown, particularly in the context of NMT. One method to probe the interpretability of neural networks is through Topological Data Analysis (TDA), where topological features that are intrinsic to the model are extracted and explained, which is also underexplored in NMT. Therefore, in this study, we propose to apply TDA to explain the power of transformer on the task of NMT. Since the attention mechanism is the core of transformer, topology-related techniques are applied to analyze the attention maps generated by transformer models during translation. We place the major focus on the French-English and Chinese-English translation tasks. The specific research question is as follows:

Do French and English sentences create similar or different topological structures in the attention maps generated by transformer models during translation? If so, do topological differences in attention maps correlate with translation quality? What about Chinese and English sentences?

2. Background

2.1. Algebraic Topology

Topology is a branch of mathematics that characterizes shapes, spaces, and sets by their connectivity (Guss & Salakhutdinov, 2018). In particular, algebraic topology is a subfield of topology that assigns algebraic properties, such as groups and chains, to topological spaces to enable more expressive characterization and analysis. Formally, let X be a compact metric space. We can define a p -simplex as a p -dimensional object determined by $p + 1$ vertices $\{x_0, \dots, x_p\} \subseteq X$. Depending on the value of p , these simplices bear different names:

- $p = 0$: point
- $p = 1$: line segment
- $p = 2$: triangle
- $p = 3$: tetrahedron
- ...

Now consider a collection of such p -simplices, called \mathcal{K} . Then \mathcal{K} is called a *simplicial complex* if it satisfies these conditions:

1. If $\sigma \in \mathcal{K}$ and τ is a face of σ , then $\tau \in \mathcal{K}$;
2. If $\sigma_1, \sigma_2 \in \mathcal{K}$, then $\sigma_1 \cap \sigma_2 = \emptyset$ or $\sigma_1 \cap \sigma_2 \in \mathcal{K}$.

Given a simplicial complex \mathcal{K} in the compact metric space X , one method that is frequently used to study \mathcal{K} is homology. The core idea of homology is to construct chains, cycles, and boundaries from the simplices in \mathcal{K} and analyze their relationships. Given a dimension n , the n th homology group of the simplicial complex \mathcal{K} is defined as $H_n(\mathcal{K}) = \mathbb{Z}^{\beta_n}$, where β_n is called the *n th Betti number*. For $n \geq 1$, the n th Betti number β_n measures the number of n -dimensional holes in the simplicial complex \mathcal{K} , while β_0 measures the number of connected components in \mathcal{K} . For example, a torus is a 2-dimensional surface with 1 connected component, 2 1-dimensional holes, and 1 2-dimensional void. Therefore, the homology groups of the torus are $H_0(\mathcal{K}) = \mathbb{Z}^1$, $H_1(\mathcal{K}) = \mathbb{Z}^2$, and $H_2(\mathcal{K}) = \mathbb{Z}^1$.

2.2. Persistent Homology

Realistically, given a collection of points in \mathbb{R}^n , we would like to extract meaningful topological information that characterizes these points. Persistent homology is one method that computes topological characteristics of such point collections. Given a collection of points P , we first construct a simplicial complex \mathcal{K} known as the Vietoris-Rips (VR) complex. The vertices in \mathcal{K} are the points in P . To build the edges that connect the points, we consider an increasing sequence of radii. For each radius r in the sequence, we superimpose a ball of radius r centered at each point in P . We connect two points x_1 and x_2 with an edge when a ball centered at one point first covers the other, which occurs when the distance between them satisfies $d(x_1, x_2) \leq r$.

Notice that as the radius r increases, more and more edges are connected, leading to the emergence and disappearance of topological features. We thus can characterize these topological features by their emergence time and disappearance time, or birth time and death time, using standard topology terminology. For instance, recall our previous example concerning points x_1 and x_2 . If a 1-dimensional hole α emerges because of the addition of the edge between them at r_1 , then we would say that α has a birth time of r_1 . Similarly, if at some later radius $r_2 > r_1$ the hole α disappears because of the addition of another edge, then we would denote r_2 as the death time of α .

Figure 1 below shows a visualization of computing persistent homology using the method described above on a set of points in \mathbb{R}^2 . The left figure shows the sequence of radii increasing from 0 to 6.15, along which edges are added to the simplicial complex. Note that at $r = 5.6$, the addition of

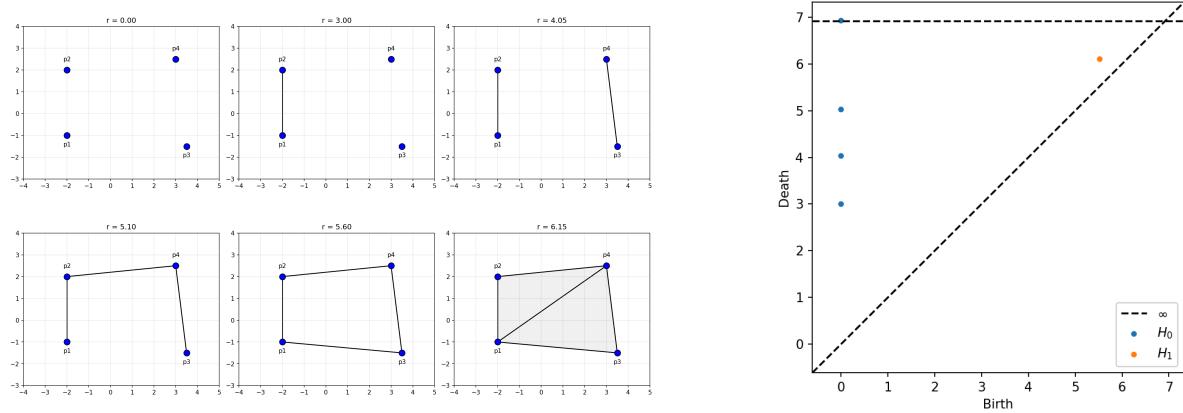


Figure 1: Visualization of Vietoris-Rips filtration on a set of points in \mathbb{R}^2 .

the edge between p_1 and p_3 forms a quadrilateral in the simplicial complex, which results in a 1-dimensional hole. Subsequently, at $r = 6.15$, the hole disappears because of the addition of the edge between p_1 and p_4 . The simplicial complex constructed by increasing the radius r is called a *filtration*. On the other hand, the right figure shows the persistence diagram of this filtration. Note that the 1-dimensional feature described previously corresponds to the orange point at $(5.6, 6.15)$ in the persistence diagram.

Now, given two sets of points, we can compute their persistence diagrams using persistent homology, but we need a metric to compare the two persistence diagrams. *Wasserstein distance* is the most common metric for this task. Given persistence diagrams D_1 and D_2 , Wasserstein distance finds the best matching between points in D_1 and D_2 and computes the sum of distances between matched points. Formally, let p be a fixed dimension. The p -Wasserstein distance between D_1 and D_2 is defined as:

$$W_p(D_1, D_2) = \inf_{\phi: D_1 \rightarrow D_2} \left(\sum_{x \in D_1} \|x - \phi(x)\|^p \right)^{1/p}.$$

In practice, we compute Wasserstein distances separately for each homological dimension and sum them to obtain the total distance:

$$W(D_1, D_2) = \sum_{n=0}^{\infty} W_n(D_1, D_2).$$

3. Related Work

There are numerous past studies that focused on using algebraic topology to analyze neural networks. For example, the paper by Bianchini & Scarselli (2014) is a pioneer study that leveraged topological tools to compare the expressivity of shallow and deep neural networks. They discovered that for deep neural networks, the sum of the Betti numbers that the network can express can grow exponentially with the number of hidden units. Later, Guss & Salakhutdinov (2018) extended this work by empirically applying Betti numbers to measure the topological complexity of real-world datasets and characterize the expressivity of fully-connected neural networks. These studies laid a solid foundation for using topological tools to study neural networks, but the generalization of such techniques to more complex neural structures still remains limited.

In addition to studying neural networks using topology, there are also attempts to apply topological techniques to language modeling. Fitz (2022) introduces the notion of a *word manifold*, which turns n -gram models on raw texts of various languages into simplicial complexes, allowing for topological analysis. This study is foundational in applying TDA to NLP tasks, but it lacks mention of neural networks. More recently, Draganov & Skiena (2024) makes an effort to study word embeddings generated by large language models by considering the d -dimensional space that these embeddings are located in as a topological space. Then, they applied persistent homology to extract topological patterns from the embedding spaces formed by 81 languages. Their study suggested statistically significant results that word embeddings carry meaningful linguistic information, but there was no analysis of the underlying neural models either.

Perhaps the most closely related paper to our study is the one by Meirom & Bobrowski (2022). In this study, they also looked at embeddings as Draganov & Skiena (2024) did, but Meirom & Bobrowski (2022) argue that certain semantics are inherent to the real world and are not language dependent. For example, *dog* and *cat* are both common pets, so they often appear in the same context regardless of the language. Under this assumption, they claimed that the embedding spaces of different languages should be isomorphic to each other at the sentence level, and their results supported this. An interesting question therefore arose: since the embedding spaces of different languages at the sentence level are isomorphic, and an NMT system transforms a sentence from the source language to the target language, how does the NMT system preserve such isomorphism during translation? This question is thus the main motivation of the current study.

Some previous studies also attempted to interpret transformers by analyzing attention. For example, Ravishankar et al. (2021) studied fully using the attention of multilingual BERT to decode syntactic

dependency trees of 18 languages, including French and English, and their results showed that solely using attention can achieve competitive accuracy in dependency parsing, suggesting that attention does encode meaningful syntactic information, which could be helpful in translation as well. Furthermore, Kushnareva et al. (2021) studied the attention mechanism with topology. In the study, they first built weighted graphs from attention maps by treating tokens as nodes and attention weights as edges, followed by applying persistent homology on the graph to construct a filtration. Their topic was on detecting artificially generated texts, which is different from our study, but the process was inspiring for our methodology.

To conclude this section, Uchendu & Le (2025) in their paper of a survey on using TDA to approach NLP problems stated that:

One glaring application is on multi-lingual tasks... Due to the benefits of TDA which include performing robustly on heterogeneous, imbalanced, and noisy data, its application on multi-lingual tasks is necessary.

Hence, the current study of applying TDA to NMT, which is a multi-lingual task, is motivated.

4. Methodology

This section presents the methodology used in this study. Section 4.1 gives an overview of the model selected in the study. Then, Section 4.2 introduces the datasets from which the sentences for analysis are drawn. This section closes with Section 4.3 that details the experimental design and analysis procedure.

4.1. Model

This study revolves around studying the attention mechanisms of NMT systems. Therefore, a pre-trained NMT system must be selected so that we can extract the attention maps to analyze. In this study, the NLLB (No Language Left Behind) model developed by Meta is chosen (Team et al., 2022). This model offers translation between 200 languages, including English, French, and Chinese, and is open-sourced. The NLLB models are available on Hugging Face, ranging from 600M to 54B parameters. The distilled NLLB with 1.3B parameters is selected for its balance between performance and computational cost. This model has 24 encoder layers and 24 decoder layers, each layer containing 16 attention heads.

4.2. Datasets

The selected model must be run on some sentences for translation to generate attention maps. The evaluation datasets curated by the WMT are picked. WMT is a well-known NMT benchmark that hosts annual MT competitions, each year with a different combination of source and target languages. In this study, the WMT14 benchmark is chosen for the French-English analysis, and the 2017 WMT (WMT17) benchmark is chosen for the Chinese-English analysis due to their popularity. As mentioned before, Vaswani et al. (2017) also used WMT14 to evaluate their transformer model, which proves the suitability of this dataset for NMT research.

The WMT14 French-English benchmark contains 3,000 validation sentence pairs, while the WMT17 Chinese-English benchmark contains only 2,000 validation sentence pairs. To ensure a fair comparison between languages, we selected only the first 2,000 sentence pairs from the WMT14 French-English validation dataset and all 2,000 sentence pairs from the WMT17 Chinese-English validation dataset to conduct the experiment.

4.3. Experiment Design

We would like to analyze whether the attention maps generated by the NMT model are different for each language. Therefore, for each French-English sentence pair in the datasets, the model is run on the English sentence to generate the French translation, as well as on the French sentence to generate the English translation. This way, the encoder of the model processes the same sentence twice, but in two different languages, so we can extract the encoder attention maps for both languages. The experiment is repeated for the Chinese-English sentence pairs. Upon generating the translations, only the attention maps in the last layer of the encoder are extracted for analysis, since the last layer is expected to contain the most refined attention information. The attention weights across the 16 heads are mean-aggregated to form a single attention map for each sentence.

After the attention maps for all the sentence pairs are extracted, we apply topological methods to analyze these attention maps. For each attention map, we build a weighted graph by treating the words in the sentence as nodes and adapting the attention weights as edges, following Kushnareva et al. (2021). The weight of each edge has the value of $1 - \alpha$, where α is the attention weight between two words. This way, a higher attention weight would correspond to a smaller edge weight, resembling a shorter distance between two nodes. The constructed graph is undirected, meaning that although attention is not symmetrical, the distance between two words would be taken as one minus the maximum attention weight between them.

Upon building the weighted graph for each attention map, we run the filtration process described in Section 2.2 to compute the persistence diagram using persistent homology. Consequently, for each sentence pair, this step results in two persistence diagrams, one for the English sentence and one for the corresponding sentence in French or Chinese. In this study, only zeroth-order and first-order topological features are considered for two reasons. First, the French-English dataset has mean sentence lengths of 17.9 tokens for English and 19.4 tokens for French, while the means for the Chinese-English dataset are 26.0 for English and 43.9 for Chinese. These numbers indicate that the VR complexes for the sentence pairs are relatively small, making higher-order topological features less likely to appear. Second, zeroth-order and first-order features are easier to interpret in the context of attention maps than higher-order features, which may not have clear meanings in this setting.

The corresponding computed persistence diagrams for each sentence pair permit the calculation of the Wasserstein distance between the two diagrams, from which we learn about how the NLLB model attends to different languages at a topological level. In this context, a smaller Wasserstein distance indicates that the attention maps for the two languages are topologically similar, which shows that the system processes the two languages in a similar manner. Conversely, a larger Wasserstein distance indicates that the attention maps are topologically different, suggesting that the model treats the two languages differently.

Lastly, we would like to see if topological differences in attention maps can be an indication of translation quality. Therefore, given the translation generated in a previous step of the experiment, we compute the BLEU score by comparing the translation to the reference sentence in the dataset. For each language direction, we aggregate the BLEU scores and conduct correlation analysis with the Wasserstein distances computed previously. For this step, we hypothesize a strong negative correlation between BLEU scores and Wasserstein distances, as it is intuitive that better translations should correspond to more similar attention maps between the two languages.

5. Results & Discussion

This section presents the topological findings from the experiment described above. Section 5.1 shows the analysis of attention maps using Wasserstein distance, delving into how NLLB attends to sentences of different languages. Next, Section 5.2 presents some insights into the translation quality of NLLB, followed by Section 5.3, which combines the results from Section 5.1 and Section 5.2 and conducts correlation analysis between topological differences and translation quality. Lastly, Section 5.4 concludes this section by discussing some limitations that hinder the model's performance on Chinese translation, provided as possible explanations for the results in Section 5.2 and Section 5.3.

5.1. Topology

Before analyzing the topological features in the attention maps, we first examine a typical attention map. The left plot of Figure 2 below shows the average self-attention in the last layer of the encoder for an English sentence. From the plot, note that the NLLB model creates a language tag at the beginning of the sentence, as well as an end-of-sentence (EOS) tag at the end of the sentence. Further notice that the EOS column of the plot has a very light color, meaning that every token in the sentence seems to attend strongly to the EOS token. This phenomenon suggests that the model considers the EOS token to be very important, perhaps using it as a scratchpad to store information about the entire sentence. The pattern is also present for the language tag, but not as strongly as EOS. However, since we are analyzing how the NMT model understands the sentence instead of byproducts of the translation process, these special tags are removed, and the remaining attention weights are renormalized to 1 for the topological analysis. The middle plot of Figure 2 shows the attention map after filtering and renormalization. Additionally, the right plot of Figure 2 shows the distance matrix of the sentence, where each distance has the value of $1 - \alpha$, with α being the attention weight between two words. The distance of every word to itself is set to 0.

With the distance matrix, the persistence diagram of this attention map can be computed. Figure 3 below shows the persistence diagrams of the same sentence presented in Figure 2 in English and French. Note that both persistence diagrams show various zeroth-order topological features, which is expected because different parts of the sentence attend differently to each other. However, first-order topological features seem to be rare and ephemeral in both diagrams. These patterns are consistent across most sentences, both in the French-English and Chinese-English datasets.

Table 1 below shows the summary statistics for the metrics computed for both the French-English and Chinese-English datasets. The table presents the minimum, maximum, mean, and median values

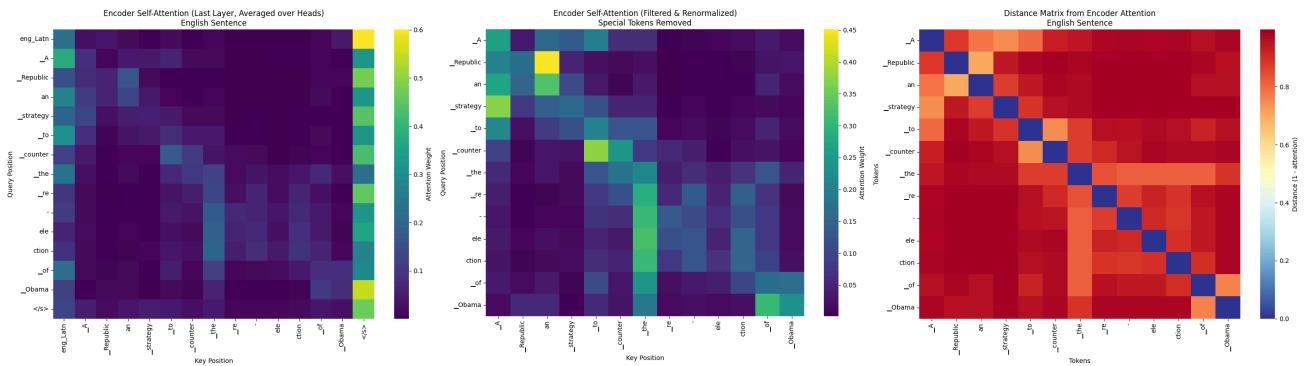


Figure 2: Attention maps and distance matrix for the English sentence “A Republican strategy to counter the re-election of Obama”.

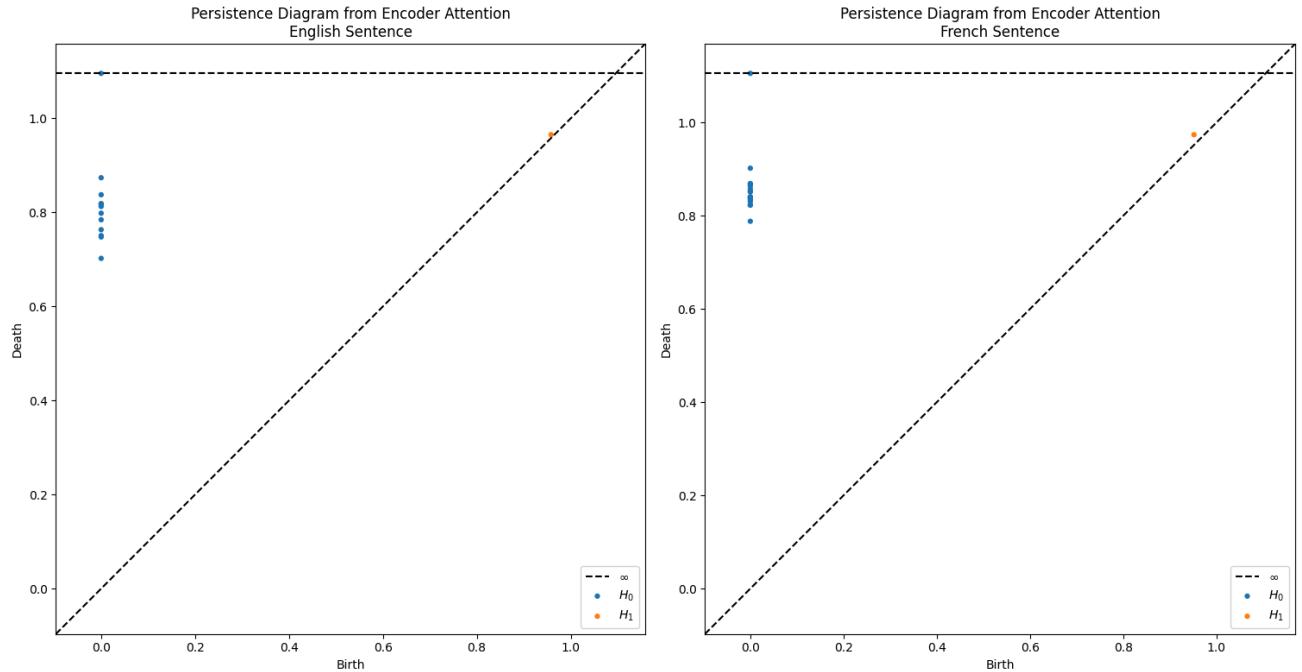


Figure 3: Persistence diagrams for the English sentence “A Republican strategy to counter the re-election of Obama” (left) and its French translation “Une stratégie républicaine pour contrer la réélection d’Obama” (right).

for Wasserstein distances, token counts, and topological features across the 2,000 sentence pairs for each language pair. From the table, note that the majority of topological differences between languages seem to stem from zeroth-order topological features, as first-order topological features only contribute to smaller than 1.0 Wasserstein distance on average for both datasets. This discovery aligns with the observation shown in Figure 3 that first-order topological features are rare and ephemeral in attention maps. Furthermore, the French sentences in our dataset are generally longer than their English counterparts, but the Chinese sentences have about the same number of tokens as their English counterparts. Nevertheless, the model still generates more H_1 features for Chinese sentences than for English sentences on average, indicating that the model attends to Chinese sentences in a more complex manner. Lastly, the last two rows of Table 1 for each language pair show the cross-language correlation of topological features. The strong correlation coefficients (> 0.65) show that NLLB creates topologically similar attention maps regardless of languages, which corroborates the statement by Meiron & Bobrowski (2022) that isomorphisms between topological structures of different languages likely exist.

Metric	Min	Max	Mean	Median
French-English				
Wasserstein Distance (Total)	0.0	40.9	5.0	4.1
Wasserstein Distance (H_0)	0.0	40.3	5.0	4.1
Wasserstein Distance (H_1)	0.0	0.5	0.1	0.0
Token Count (H_0 Features) (English)	1	117	24.7	22
Token Count (H_0 Features) (French)	2	177	31.7	28
H_1 Features (English)	0	51	4.6	3
H_1 Features (French)	0	78	5.7	4
H_0 Cross-Language Correlation			0.95	
H_1 Cross-Language Correlation			0.86	
Chinese-English				
Wasserstein Distance (Total)	0.2	26.6	4.4	3.5
Wasserstein Distance (H_0)	0.2	26.2	4.2	3.3
Wasserstein Distance (H_1)	0.0	0.7	0.1	0.1
Token Count (H_0 Features) (English)	2	109	36.0	35
Token Count (H_0 Features) (Chinese)	4	111	36.2	35
H_1 Features (English)	0	41	8.4	8
H_1 Features (Chinese)	0	67	12.6	11
H_0 Cross-Language Correlation			0.85	
H_1 Cross-Language Correlation			0.68	

Table 1: Summary statistics for topological metrics computed on the French-English and Chinese-English datasets (2,000 sentence pairs each).

5.2. Translation Quality

Now we examine the translation quality of the NLLB model on the datasets with BLEU scores. The BLEU metric is a widely used measurement of translation quality that compares the generated translation to a reference sentence. The BLEU score is comprised of a brevity penalty factor that penalizes the translation for being too short, as well as a geometric mean of modified n -gram precisions that measures how many n -grams in the translation appear in the reference sentence. The final BLEU score ranges from 0 to 100, with higher scores indicating better translation quality (Papineni et al., 2002).

Figure 4 below shows the distributions of BLEU scores for translations in the French-English and Chinese-English datasets. From the plots, we note that the Chinese-English language pair achieves lower BLEU scores on average than French-English by approximately 10 points. In particular,

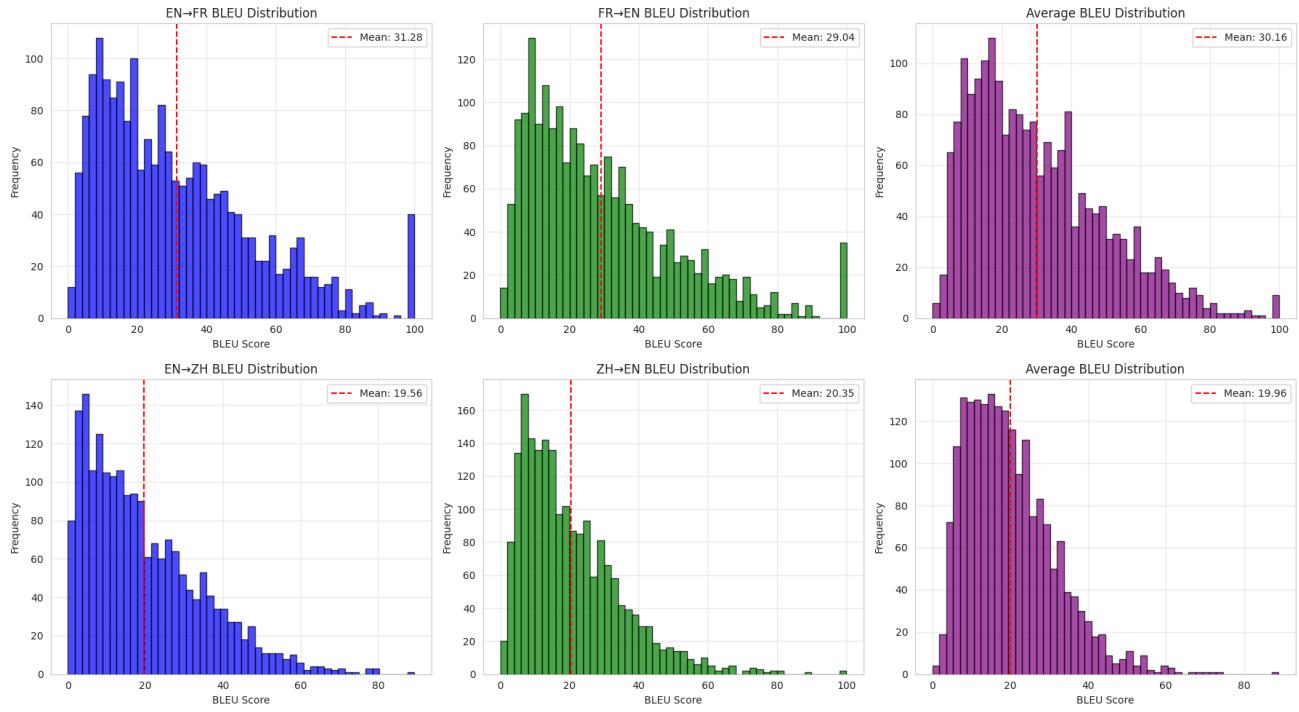


Figure 4: Distributions of BLEU scores for translations in the French-English (top) and Chinese-English (bottom) datasets.

perfect translation ($\text{BLEU} = 100$) is not uncommon between French and English but is rarely seen in Chinese-English translations. This phenomenon is expected because French and English come from the same typological roots, while Chinese belongs to a completely different language family. Therefore, it is generally more difficult to translate between Chinese and English than between French and English. A more specific analysis of translation errors is presented in Section 5.4.

5.3. Correlation Analysis

After computing the Wasserstein distances and BLEU scores for all sentence pairs in the datasets, we conduct a correlation analysis to examine whether Wasserstein distance is an indication of translation quality. Figure 5 below shows the scatter plots of Wasserstein distances and BLEU scores for the datasets. We notice that in all 6 plots, the Pearson correlation coefficients are all negative, but the magnitudes are all smaller than 0.1. All the coefficients are statistically significant at the significance level of 0.05. These results suggest that, although the correlations are weak, there is significant evidence that there are negative correlations between Wasserstein distances and BLEU scores in all language directions. This finding partially supports our earlier hypothesis that there is a negative correlation between topological differences and translation quality, but the correlation is not very strong.

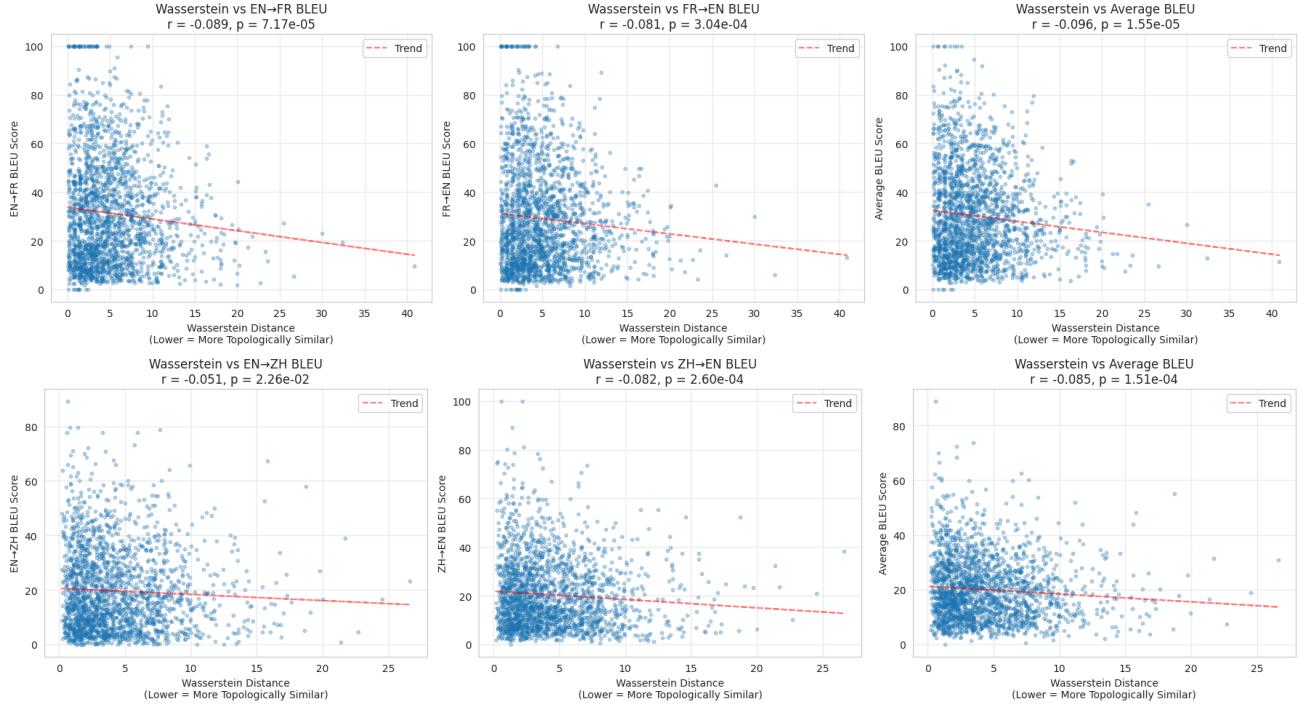


Figure 5: Scatter plots showing the relationship between Wasserstein distances and BLEU scores for the French-English (top) and Chinese-English (bottom) datasets.

Now, as shown in Table 1 above, token count directly reflects the number of zeroth-order topological features in the sentence, which can possibly confound our correlation analysis. This factor is also intuitively worrisome because longer sentences are intrinsically more difficult to translate correctly. Therefore, a correlation analysis that controls for the effect of token count is necessary. Figure 6 below shows the scatter plots for the partial correlation between Wasserstein distances and BLEU scores. Given two variables of interest for correlation analysis and one or more possibly confounding variables, partial correlation first fits two linear regression models that use the confounding variables to predict each variable of interest separately. With the linear regression models come the residuals of the two variables of interest that the confounding variables cannot explain. Then, we carry out the correlation analysis on the two sets of residuals, attempting to find correlation between the two variables of interest in the parts that are not affected by the confounding variables. In the context of this paper, the two variables of interest are the Wasserstein distances and the BLEU scores, and the confounding variable is the token count.

From Figure 6, we note that the partial correlations for the French-English pair are higher than the correlations from Figure 5 in both language directions, and the partial correlation for the average case is the highest. These results are also strongly statistically significant with very small p -values. This suggests that, even though French and English sentences have very different token counts as

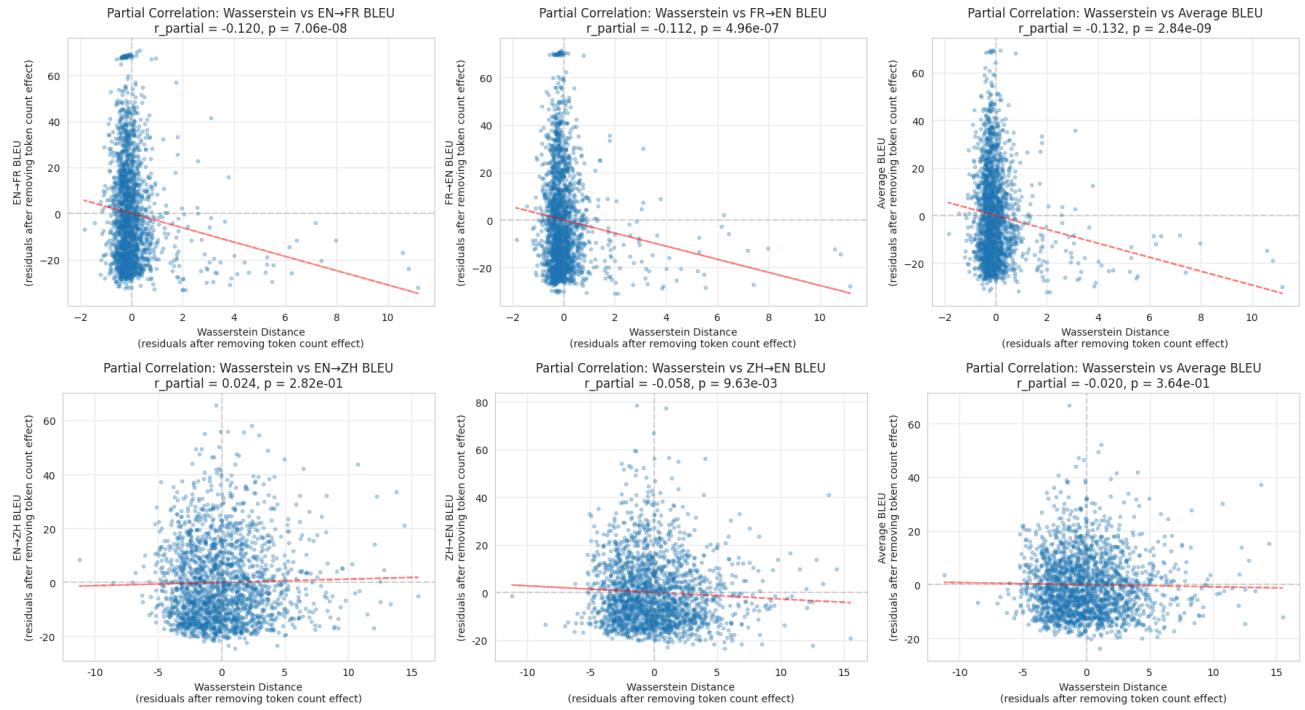


Figure 6: Scatter plots showing the relationship between Wasserstein distances and BLEU scores after controlling for token counts for the French-English (top) and Chinese-English (bottom) datasets.

indicated in Table 1, this does not seem to confound the correlation between topological differences in the attention maps and translation quality. In fact, controlling for token count renders the correlation stronger, which shows that token count is actually masking the true correlation between Wasserstein distances and BLEU scores. Therefore, we can conclude that preserving topological features in attention maps independently contributes to translation quality between French and English.

On the other hand, the partial correlations for Chinese-English become weaker as shown in Figure 6, with statistically insignificant p -values in the English-Chinese direction and the average case. This shows that token count, in the Chinese-English case, is strongly confounding the correlation between Wasserstein distances and BLEU scores. In other words, the difficulty in translating longer sentences stands out more in this case, possibly overwriting the topological differences in the attention maps that the model generates. This provides a clear contrast to the French-English analysis. In Section 5.4 below, we will discuss possible limitations that may explain this phenomenon.

5.4. Model Limitations: Chinese Translation Truncation

The Chinese-English results reveal an important limitation of the NLLB model that affects the reliability of our correlation analysis. As noted in Section 5.2, Chinese translations achieve approximately 10 points lower BLEU scores than French translations on average. To understand

whether this represents systematic model failure or expected typological difficulty, we analyzed 10 Chinese-English sentence pairs with the lowest average BLEU scores. Upon investigation, we discovered that the NLLB model frequently truncates English-to-Chinese translations. For example, the English sentence

At 10:00pm, Sun Yijie, who had been pregnant for four months, was released on bail of NT\$200,000.

is translated into Chinese as

苏伊杰已经怀孕四个月,

(*Su Yijie had been pregnant for four months,*)

The translation terminates abruptly at a comma and entirely omits the latter part of the source sentence. This is not an isolated case. In fact, this truncation issue affects 8 out of the 10 lowest-scoring sentence pairs. Further experimentation confirmed that the problem persists across multiple decoding strategies, including greedy decoding and beam search, and occurs specifically in the English-to-Chinese direction, not Chinese-to-English, nor in French-English translations.

This systematic truncation has two important implications. First, it severely impacts BLEU scores in a way that is unrelated to topological structures, which confounds the correlation analysis for Chinese-English. The partial correlation results in Section 5.3 should therefore be interpreted cautiously for this language pair. Second, it represents a systematic issue with the NLLB-1.3B model that may be relevant for other researchers working with English-to-Chinese translation using this model. Despite these complications with Chinese-English, the contrast with the French-English analysis shows that preserving topological structures significantly correlates with translation quality when translations are complete and of high quality.

6. Conclusion

To summarize, this paper attempts to contribute to interpreting NMT systems using tools from topological data analysis. A particular focus is placed on analyzing the attention maps that NMT models generate during the translation process. In the study, the NLLB NMT model with 1.3B parameters developed by Meta is selected for analysis, due to its wide inclusion of languages and power on this task (Team et al., 2022). Both French-English and Chinese-English language pairs are selected, considering their typological differences. From the WMT14 and WMT17 MT benchmarks, 2,000 French-English and Chinese-English sentence pairs are acquired separately, and

the corresponding attention maps are extracted for TDA. Lastly, the Wasserstein distances computed from persistent homology are correlated with BLEU scores, which measure translation quality.

Before the study, the hypothesis was that higher topological differences, or larger Wasserstein distances between attention maps in the encoder for different languages, are correlated with lower translation quality because the NMT model is likely to process the two languages in a similar manner and translate better if the attention maps are topologically similar. From the analysis, we demonstrate that topological dissimilarity in attention maps independently correlates with translation quality for French-English after controlling for sentence length ($r = -0.132, p = 2.84 \times 10^{-9}$). While the magnitude is modest, this finding reveals that transformers preserve meaningful topological structure across languages, which contributes meaningfully to success in translation. For Chinese-English, however, the correlation becomes negligible after controlling for sentence length ($r = -0.020, p = 0.364$), likely due to systematic truncation issues in English-to-Chinese translations by the NLLB model that confound the analysis.

6.1. Future Directions

As mentioned in Section 5.4, the problem that the NLLB model generates truncated Chinese translations from English source sentences has been discovered. Therefore, it is recommended that future researchers address this problem by either fixing the bug intrinsically or changing the model completely. After resolving this issue, the Chinese translations can hopefully be more accurate, which would allow a more reliable correlation analysis given better BLEU scores. In addition, this study only looks at French-English and Chinese-English language pairs. Future studies are encouraged to choose other language pairs that have different typological relationships, such as German-English or Arabic-English, to see if the findings in this study generalize. Lastly, topology is only one of the numerous ways to measure differences. Future researchers are welcome to explore other methods, such as geometric approaches, to analyze attention maps in NMT systems.

7. Acknowledgements

I am extremely grateful to Professor Gerald Penn and TA Jinman Zhao for giving valuable feedback and conceptual guidance in algebraic topology throughout the process of this research.

Bibliography

- [1] A. Vaswani *et al.*, “Attention is All you Need,” in *Advances in Neural Information Processing Systems*, Curran Associates, Inc., 2017. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf
- [2] W. H. Guss and R. Salakhutdinov, “On Characterizing the Capacity of Neural Networks using Algebraic Topology,” *CoRR*, 2018, [Online]. Available: <https://arxiv.org/abs/1802.04443>
- [3] M. Bianchini and F. Scarselli, “On the Complexity of Neural Network Classifiers: A Comparison Between Shallow and Deep Architectures,” *IEEE Transactions on Neural Networks and Learning Systems*, vol. 25, no. 8, pp. 1553–1565, 2014.
- [4] S. Fitz, “The Shape of Words - topological structure in natural language data ,” in *Proceedings of Topological, Algebraic, and Geometric Learning Workshops 2022*, in Proceedings of Machine Learning Research, vol. 196. PMLR, 2022, pp. 116–123.
- [5] O. Draganov and S. Skiena, “The Shape of Word Embeddings: Quantifying Non-Isometry with Topological Data Analysis,” in *Findings of the Association for Computational Linguistics: EMNLP 2024*, Miami, Florida, USA: Association for Computational Linguistics, Nov. 2024, pp. 12080–12099.
- [6] S. H. Meirom and O. Bobrowski, “Unsupervised Geometric and Topological Approaches for Cross-Lingual Sentence Representation and Comparison,” in *Proceedings of the 7th Workshop on Representation Learning for NLP*, Dublin, Ireland: Association for Computational Linguistics, May 2022, pp. 173–183.
- [7] V. Ravishankar, A. Kulmizev, M. Abdou, A. Søgaard, and J. Nivre, “Attention Can Reflect Syntactic Structure (If You Let It),” in *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, Online: Association for Computational Linguistics, Apr. 2021, pp. 3031–3045.
- [8] L. Kushnareva *et al.*, “Artificial Text Detection via Examining the Topology of Attention Maps,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Online and Punta Cana, Dominican Republic: Association for Computational Linguistics, Nov. 2021, pp. 635–649.

-
- [9] A. Uchendu and T. Le, “Unveiling Topological Structures from Language: A Comprehensive Survey of Topological Data Analysis Applications in NLP.” [Online]. Available: <https://arxiv.org/abs/2411.10298>
 - [10] N. Team *et al.*, “No Language Left Behind: Scaling Human-Centered Machine Translation .” [Online]. Available: <https://arxiv.org/abs/2207.04672>
 - [11] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, “Bleu: a Method for Automatic Evaluation of Machine Translation,” in *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, Philadelphia, Pennsylvania, USA: Association for Computational Linguistics, July 2002, pp. 311–318. doi: [10.3115/1073083.1073135](https://doi.org/10.3115/1073083.1073135).